

Systems I

Computer Systems I

Class 8

Systems I

Today

- Arrays
 - One-dimensional
 - Multi-dimensional (nested)
 - Multi-level

Systems I

Basic Data Types

- Integral
 - Stored & operated on in general (integer) registers
 - Signed vs. unsigned depends on instructions used

Intel	GAS	Bytes	C
byte	b	1	[unsigned] char
word	w	2	[unsigned] short
double word	l	4	[unsigned] int
quad word	q	8	[unsigned] long int (x86-64)


- Floating Point
 - Stored & operated on in floating point registers

Intel	GAS	Bytes	C
Single	s	4	float
Double	l	8	double
Extended	t	10/12/16	long double

Array Allocation

Basic Principle


- T $A[L]$;
 - Array of data type T and length L
 - Contiguously allocated region of $L * \text{sizeof}(T)$ bytes



Array Access

Basic Principle

- T $A[L]$;
 - Array of data type T and length L
 - Identifier A can be used as a pointer to array element 0: Type T^*




Reference Type Value

<code>val[4]</code>	
<code>val</code>	
<code>val+1</code>	
<code>&val[2]</code>	
<code>val[5]</code>	
<code>*(val+1)</code>	
<code>val + i</code>	

Array Example

```
typedef int zip_dig[5];

zip_dig cmu = { 1, 5, 2, 1, 3 };
zip_dig mit = { 0, 2, 1, 3, 9 };
zip_dig ucb = { 9, 4, 7, 2, 0 };
```



- Declaration "`zip_dig cmu`" equivalent to "`int cmu[5]`"
- Example arrays were allocated in successive 20 byte blocks
 - Not guaranteed to happen in general

Array Accessing Example

zip_dig cmu:

1	5	2	1	3
---	---	---	---	---

16 20 24 28 32 36

```
int get_digit
(zip_dig z, int dig)
{
    return z[dig];
}
```

IA32

```
# %edx = z
# %eax = dig
movl (%edx,%eax,4),%eax # z[dig]
```

- Register %edx contains starting address of array
- Register %eax contains array index
- Desired digit at $4 * \text{\%eax} + \text{\%edx}$
- Use memory reference ($\text{\%edx}, \text{\%eax}, 4$)

Referencing Examples

zip_dig cmu:

1	5	2	1	3
---	---	---	---	---

16 20 24 28 32 36

zip_dig mit:

0	2	1	3	9
---	---	---	---	---

36 40 44 48 52 56

zip_dig ucb:

9	4	7	2	0
---	---	---	---	---

56 60 64 68 72 76

Reference	Address	Value	Guaranteed?
mit[3]			
mit[5]			
mit[-1]			
cmu[15]			

Referencing Examples

zip_dig cmu:

1	5	2	1	3
---	---	---	---	---

16 20 24 28 32 36

zip_dig mit:

0	2	1	3	9
---	---	---	---	---

36 40 44 48 52 56

zip_dig ucb:

9	4	7	2	0
---	---	---	---	---

56 60 64 68 72 76

Reference	Address	Value	Guaranteed?
mit[3]	$36 + 4 * 3 = 48$	3	Yes
mit[5]	$36 + 4 * 5 = 56$	9	No
mit[-1]	$36 + 4 * -1 = 32$	3	No
cmu[15]	$16 + 4 * 15 = 76$??	No

- No bound checking
- Out of range behavior implementation-dependent
- No guaranteed relative allocation of different arrays

Array Loop Example

Original

```
int zd2int(zip_dig z)
{
    int i;
    int zi = 0;
    for (i = 0; i < 5; i++) {
        zi = 10 * zi + z[i];
    }
    return zi;
}
```

Transformed

- As generated by GCC
- Eliminate loop variable *i*
- Convert array code to pointer code
- Express in do-while form (no test at entrance)

```
int zd2int(zip_dig z)
{
    int zi = 0;
    int *zend = z + 4;
    do {
        zi = 10 * zi + *z;
        z++;
    } while (z <= zend);
    return zi;
}
```

Array Loop Implementation (IA32)

```
int zd2int(zip_dig z)
{
    int zi = 0;
    int *zend = z + 4;
    do {
        zi = 10 * zi + *z;
        z++;
    } while(z <= zend);
    return zi;
}
```

```
# %ecx = z
xorl %eax,%eax
leal 16(%ecx),%ebx
.L59:
leal (%eax,%eax,4),%edx
movl (%ecx),%eax
addl $4,%ecx
leal (%eax,%edx,2),%eax
cmpl %ebx,%ecx
jle .L59
```

Array Loop Implementation (IA32)

Registers

%ecx z
%eax zi
%ebx zend

Computations


- $10 * zi + *z$ implemented as $*z + 2 * (zi + 4 * zi)$
- `z++` increments by 4

```
int zd2int(zip_dig z)
{
    int zi = 0;
    int *zend = z + 4;
    do {
        zi = 10 * zi + *z;
        z++;
    } while(z <= zend);
    return zi;
}
```

```
# %ecx = z
xorl %eax,%eax      # zi = 0
leal 16(%ecx),%ebx   # zend = z+4
.L59:
leal (%eax,%eax,4),%edx # 5*zi
movl (%ecx),%eax      # *z
addl $4,%ecx          # z++
leal (%eax,%edx,2),%eax # zi = *z + 2*(5*zi)
cmpl %ebx,%ecx        # z : zend
jle .L59             # if <= goto loop
```

Nested Array Example

```
#define PCOUNT 4
zip_dig pgh[PCOUNT] =
{{1, 5, 2, 0, 6},
 {1, 5, 2, 1, 3},
 {1, 5, 2, 1, 7},
 {1, 5, 2, 2, 1}};
```




- “zip_dig pgh[4]” equivalent to “int pgh[4][5]”
 - Variable **pgh**: array of 4 elements, allocated contiguously
 - Each element is an array of 5 **int**'s, allocated contiguously
- “Row-Major” ordering of all elements guaranteed

Multidimensional (Nested) Arrays

- **Declaration**

```
T A[R][C];
```


 - 2D array of data type *T*
 - *R* rows, *C* columns
 - Type *T* element requires *K* bytes



- **Array Size**
 - $R * C * K$ bytes

- **Arrangement**
 - Row-Major Ordering


```
int A[R][C];
```



Nested Array Row Access

- **Row Vectors**
 - **A[i]** is array of *C* elements
 - Each element of type *T* requires *K* bytes
 - Starting address **A** + $i * (C * K)$

```
int A[R][C];
```



Nested Array Row Access Code

```
int *get_pgh_zip(int index)
{
    return pgh[index];
}
```

```
#define PCOUNT 4
zip_dig pgh[PCOUNT] =
{{1, 5, 2, 0, 6},
 {1, 5, 2, 1, 3},
 {1, 5, 2, 1, 7},
 {1, 5, 2, 2, 1}};
```

- What data type is `pgh[index]`?
- What is its starting address?

```
# %eax = index
leal (%eax,%eax,4),%eax
leal pgh(,%eax,4),%eax
```

Will disappear
Blackboard?

Nested Array Row Access Code

```
int *get_pgh_zip(int index)
{
    return pgh[index];
}
```

```
#define PCOUNT 4
zip_dig pgh[PCOUNT] =
{{1, 5, 2, 0, 6},
 {1, 5, 2, 1, 3},
 {1, 5, 2, 1, 7},
 {1, 5, 2, 2, 1}};
```


```
# %eax = index
leal (%eax,%eax,4),%eax # 5 * index
leal pgh(,%eax,4),%eax # pgh + (20 * index)
```

- Row Vector
 - `pgh[index]` is array of 5 int's
 - Starting address `pgh+20*index`
- IA32 Code
 - Computes and returns address
 - Compute as `pgh + 4*(index+4*index)`

Nested Array Row Access

- Array Elements
 - `A[i][j]` is element of type `T`, which requires `K` bytes
 - Address $A + i * (C * K) + j * K = A + (i * C + j) * K$

```
int A[R][C];
```



Two dimensional array “facts”

- `slot_type A[R][C];`
- Let K = type size (4 for int, 8 for double, 1 for char, etc.)
- Let R = 1st dimension (number of rows)
- Let C = 2nd dimension (number of columns)
- Size of row = $K * C$ (cell size * number of cells per row)
- Size of A = $(K * C) * R$ (row size * number of rows)
- A is a pointer to the beginning of the whole table
- $A[i]$ is a pointer to the beginning of row i = $A + (i * K * C)$
- $A[i][j]$ is the entry in row i , column j . It can be found at address $A + (i * K * C) + j * K$
- What is the value of $A+1$?


Nested Array Element Access Code

```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```

```
# %ecx = dig
# %eax = index
leal 0(,%ecx,4),%edx    # 4*dig
leal (%eax,%eax,4),%eax  # 5*index
movl pgh(%edx,%eax,4),%eax # *(pgh + 4*dig + 20*index)
```


- **Array Elements**
 - `pgh[index][dig]` is int
 - Address: `pgh + 20*index + 4*dig`
- **IA32 Code**
 - Computes address `pgh + 4*dig + 4*(index+4*index)`
 - `movl` performs memory reference

Strange Referencing Examples



Reference	Address	Value	Guaranteed?
<code>pgh[3][3]</code>	Will disappear		
<code>pgh[2][5]</code>			
<code>pgh[2][-1]</code>			
<code>pgh[4][-1]</code>			
<code>pgh[0][19]</code>			
<code>pgh[0][-1]</code>			

Strange Referencing Examples



Reference	Address	Value	Guaranteed?
pgh[3][3]	$76 + 20 \cdot 3 + 4 \cdot 3 = 148$	2	Yes
pgh[2][5]	$76 + 20 \cdot 2 + 4 \cdot 5 = 136$	1	Yes
pgh[2][-1]	$76 + 20 \cdot 2 + 4 \cdot (-1) = 112$	3	Yes
pgh[4][-1]	$76 + 20 \cdot 4 + 4 \cdot (-1) = 152$	1	Yes
pgh[0][19]	$76 + 20 \cdot 0 + 4 \cdot 19 = 152$	1	Yes
pgh[0][-1]	$76 + 20 \cdot 0 + 4 \cdot (-1) = 72$??	No


- Code does not do any bounds checking
- Ordering of elements within array guaranteed

Multi-Level Array Example

```
zip_dig cmu = { 1, 5, 2, 1, 3 };
zip_dig mit = { 0, 2, 1, 3, 9 };
zip_dig ucb = { 9, 4, 7, 2, 0 };
```

```
#define UCOUNT 3
int *univ[UCOUNT] = {mit, cmu, ucb};
```

- Variable `univ` denotes array of 3 elements
- Each element is a pointer
 - 4 bytes
- Each pointer points to array of int's



Element Access in Multi-Level Array

```
int get_univ_digit
(int index, int dig)
{
    return univ[index][dig];
}
```

```
# %ecx = index
# %eax = dig
leal 0(,%ecx,4),%edx
movl univ(%edx),%edx
movl (%edx,%eax,4),%eax
```

Will disappear
Blackboard?

Element Access in Multi-Level Array

```
int get_univ_digit
(int index, int dig)
{
    return univ[index][dig];
}
```

```
# %ecx = index
# %eax = dig
leal 0(,%ecx,4),%edx # 4*index
movl univ(%edx),%edx # Mem[univ+4*index]
movl (%edx,%eax,4),%eax # Mem[...+4*dig]
```


Computation (IA32)

- Element access **Mem[Mem[univ+4*index]+4*dig]**
- Must do two memory reads
 - First get pointer to row array
 - Then access element within array

Array Element Accesses


Nested array

```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```



Multi-level array

```
int get_univ_digit
(int index, int dig)
{
    return univ[index][dig];
}
```




C Access looks similar, but underlying computation is different:

`Mem[pgh+20*index+4*dig]`

`Mem[Mem[univ+4*index]+4*dig]`


Strange Referencing Examples



Reference	Address	Value	Guaranteed?
<code>univ[2][3]</code>			
<code>univ[1][5]</code>			
<code>univ[2][-1]</code>			
<code>univ[3][-1]</code>			
<code>univ[1][12]</code>			

Will disappear

Strange Referencing Examples



Reference	Address	Value	Guaranteed?
<code>univ[2][3]</code>	$56 + 4 \times 3 = 68$	2	Yes
<code>univ[1][5]</code>	$16 + 4 \times 5 = 36$	0	No
<code>univ[2][-1]</code>	$56 + 4 \times -1 = 52$	9	No
<code>univ[3][-1]</code>	??	??	No
<code>univ[1][12]</code>	$16 + 4 \times 12 = 64$	7	No

- Code does not do any bounds checking
- Ordering of elements in different arrays not guaranteed

N X N Matrix Code

Fixed dimensions

- Know value of N at compile time

```
#define N 16
typedef int fix_matrix[N][N];
/* Get element a[i][j] */
int fix_ele
(fix_matrix a, int i, int j)
{
    return a[i][j];
}
```

Variable dimensions, explicit indexing

- Traditional way to implement dynamic arrays

```
#define IDX(n, i, j) ((i)*(n)+(j))
/* Get element a[i][j] */
int vec_ele
(int n, int *a, int i, int j)
{
    return a[IDX(n,i,j)];
}
```

Variable dimensions, implicit indexing

- Now supported by gcc

```
/* Get element a[i][j] */
int var_ele
(int n, int a[n][n], int i, int j)
{
    return a[i][j];
}
```

16 X 16 Matrix Access

Array Elements

- Address $A + i * (C * K) + j * K$
- $C = 16, K = 4$

```
/* Get element a[i][j] */
int fix_ele(fix_matrix a, int i, int j) {
    return a[i][j];
}
```

```
movl 12(%ebp), %edx # i
sall $6, %edx # i*64
movl 16(%ebp), %eax # j
sall $2, %eax # j*4
addl 8(%ebp), %eax # a + j*4
movl (%eax,%edx), %eax # *(a + j*4 + i*64)
```

n X n Matrix Access

■ Array Elements

- Address $A + i * (C * K) + j * K$
- $C = n, K = 4$

```
/* Get element a[i][j] */
int var_ele(int n, int a[n][n], int i, int j) {
    return a[i][j];
}
```

```
movl 8(%ebp), %eax    # n
sall $2, %eax         # n*4
movl %eax, %edx       # n*4
imull 16(%ebp), %edx  # i*n*4
movl 20(%ebp), %eax   # j
sall $2, %eax         # j*4
addl 12(%ebp), %eax   # a + j*4
movl (%eax,%edx), %eax # *(a + j*4 + i*n*4)
```

Optimizing Fixed Array Access



• Computation

- Step through all elements in column j

• Optimization

- Retrieving successive elements from single column

```
#define N 16
typedef int fix_matrix[N][N];
```

```
/* Retrieve column j from array */
void fix_column
(fix_matrix a, int j, int *dest)
{
    int i;
    for (i = 0; i < N; i++)
        dest[i] = a[i][j];
}
```

Optimizing Fixed Array Access

• Optimization

- Compute $ajp = \&a[i][j]$
 - Initially $= a + 4*j$
 - Increment by $4*N$

Register	Value
%ecx	ajp
%ebx	dest
%edx	i

```
/* Retrieve column j from array */
void fix_column
(fix_matrix a, int j, int *dest)
{
    int i;
    for (i = 0; i < N; i++)
        dest[i] = a[i][j];
}
```

```
.L8:
movl (%ecx), %eax    # Read *ajp
movl %eax, (%ebx,%edx,4) # Save in dest[i]
addl $1, %edx        # i++
addl $64, %ecx       # ajp += 4*N
cmpl $16, %edx       # i:N
jne .L8              # if !=, goto loop
```

Optimizing Variable Array Access

- Compute `ajp = &a[i][j]`
 - Initially `= a + 4*j`
 - Increment by `4*n`

Register	Value
<code>%ecx</code>	<code>ajp</code>
<code>%edi</code>	<code>dest</code>
<code>%edx</code>	<code>i</code>
<code>%ebx</code>	<code>4*n</code>
<code>%esi</code>	<code>n</code>

```

/* Retrieve column j from array */
void var_column
(int n, int a[n][n],
 int j, int *dest)
{
    int i;
    for (i = 0; i < n; i++)
        dest[i] = a[i][j];
}

```

```

.L18:
    movl (%ecx), %eax      # Read *ajp
    movl %eax, (%edi,%edx,4) # Save in dest[i]
    addl $1, %edx          # i++
    addl $ebx, %ecx        # ajp += 4*n
    cmpl $edx, %esi        # n:i
    jg .L18               # if >, goto loop

```

Rest of Today

- Memory layout
- Buffer overflow, worms, and viruses

IA32 Linux Memory Layout

- Stack**
 - Runtime stack (8MB limit)
- Heap**
 - Dynamically allocated storage
 - When call `malloc()`, `calloc()`, `new()`
- Data**
 - Statically allocated data
 - E.g., arrays & strings declared in code
- Text**
 - Executable machine instructions
 - Read-only

not drawn to scale

Upper 2 hex digits = 8 bits of address → 08 00

Memory Allocation Example

```

char big_array[1<<24]; /* 16 MB */
char huge_array[1<<28]; /* 256 MB */

int beyond;
char *p1, *p2, *p3, *p4;

int useless() { return 0; }

int main()
{
    p1 = malloc(1 << 28); /* 256 MB */
    p2 = malloc(1 << 8); /* 256 B */
    p3 = malloc(1 << 28); /* 256 MB */
    p4 = malloc(1 << 8); /* 256 B */
    /* Some print statements ... */
}

```

Where does everything go?

not drawn to scale

IA32 Example Addresses

address range $\sim 2^{32}$

\$esp	0xffffbcd0
p3	0x65586008
p1	0x55585008
p4	0x1904a110
p2	0x1904a008
&p2	0x18049760
beyond	0x08049744
big_array	0x18049780
huge_array	0x08049760
main()	0x080483c6
useless()	0x08049744
final malloc()	0x006be166

malloc() is dynamically linked
address determined at runtime

not drawn to scale

x86-64 Example Addresses

address range $\sim 2^{47}$

\$rsp	0x7fffffff8d1f8
p3	0x2aaabaadd010
p1	0x2aaaaaad010
p4	0x000011501120
p2	0x000011501010
&p2	0x000010500a60
beyond	0x000000500a44
big_array	0x000010500a80
huge_array	0x000000500a50
main()	0x000000400510
useless()	0x000000400500
final malloc()	0x00386ae6a170

malloc() is dynamically linked
address determined at runtime

not drawn to scale

C operators

Operators	Associativity
() [] -> .	left to right
! ~ ++ -- + - * & (type) sizeof	right to left
* / %	left to right
+ -	left to right
<< >>	left to right
< <= > >=	left to right
== !=	left to right
&	left to right
^	left to right
	left to right
&&	left to right
	left to right
?:	right to left
= += -= *= /= %= &= ^= != <<= >>=	right to left
,	left to right

- -> has very high precedence
- () has very high precedence
- monadic * just below

C Pointer Declarations: Test Yourself!

int *p	p is a pointer to int
int *p[13]	
int *(p[13])	
int **p	p is a pointer to a pointer to an int
int (*p)[13]	
int *f()	f is a function returning a pointer to int
int (*f)()	f is a pointer to a function returning int
int (*(*f())[13])()	
int (*(*x[3])())[5]	x is an array[3] of pointers to functions returning pointers to array[5] of ints

C Declarations (see book and WWW)

int *p	p is a pointer to int
int *p[13]	p is an array[13] of pointer to int
int *(p[13])	p is an array[13] of pointer to int
int **p	p is a pointer to a pointer to an int
int (*p)[13]	p is a pointer to an array[13] of int
int *f()	f is a function returning a pointer to int
int (*f)()	f is a pointer to a function returning int
int (*(*f())[13])()	f is a function returning ptr to an array[13] of pointers to functions returning int
int (*(*x[3])())[5]	x is an array[3] of pointers to functions returning pointers to array[5] of ints

Avoiding Complex Declarations

- Use `typedef` to build up the declaration
- Instead of `int ((*x[3])())[5]:`


```
typedef int fiveints[5];
typedef fiveints* p5i;
typedef p5i (*pfr_p5is)();
pfr_p5is x[3];
```
- `x` is an array of 3 elements, each of which is a pointer to a function returning an array of 5 ints

Rest of Today

- Memory layout
- Buffer overflow, worms, and viruses

Internet Worm and IM War

- **November, 1988**
 - Internet Worm attacks thousands of Internet hosts.
 - How did it happen?
- **July, 1999**
 - Microsoft launches MSN Messenger (instant messaging system).
 - Messenger clients can access popular AOL Instant Messaging Service (AIM) servers



Internet Worm and IM War (cont.)

- **August 1999**
 - Mysteriously, Messenger clients can no longer access AIM servers.
 - Microsoft and AOL begin the IM war:
 - AOL changes server to disallow Messenger clients
 - Microsoft makes changes to clients to defeat AOL changes.
 - At least 13 such skirmishes.
 - How did it happen?
- **The Internet Worm and AOL/Microsoft War were both based on *stack buffer overflow* exploits!**
 - many Unix functions do not check argument sizes.
 - allows target buffers to overflow.

String Library Code

Implementation of Unix function `gets()`

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

- No way to specify limit on number of characters to read
- **Similar problems with other Unix functions**
 - **`strcpy`**: Copies string of arbitrary length
 - **`scanf`, `fscanf`, `sscanf`**, when given **`%s`** conversion specification

Vulnerable Buffer Code

```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
int main()
{
    printf("Type a string:");
    echo();
    return 0;
}
```

```
unix> ./bufdemo
Type a string:1234567
1234567
```

```
unix> ./bufdemo
Type a string:12345678
Segmentation Fault
```

```
unix> ./bufdemo
Type a string:123456789ABC
Segmentation Fault
```

System 1

Buffer Overflow Disassembly

080484f0:	<echo>:		
80484f0:	55	push	%ebp
80484f1:	89 e5	mov	%esp,%ebp
80484f3:	53	push	%ebx
80484f4:	8d 5d f8	lea	0xffffffff(%ebp),%ebx
80484f7:	83 ec 14	sub	\$0x14,%esp
80484fa:	89 1c 24	mov	%ebx,(%esp)
80484fd:	e8 ae ff ff ff	call	80484b0 <gets>
8048502:	89 1c 24	mov	%ebx,(%esp)
8048505:	e8 8a fe ff ff	call	8048394 <puts@plt>
804850a:	83 c4 14	add	\$0x14,%esp
804850d:	5b	pop	%ebx
804850e:	c9	leave	
804850f:	c3	ret	
80485f2:	e8 f9 fe ff ff	call	80484f0 <echo>
80485f7:	8b 5d fc	mov	0xffffffff(%ebp),%ebx
80485fa:	c9	leave	
80485fb:	31 c0	xor	%eax,%eax
80485fd:	c3	ret	

System 1

Buffer Overflow Stack

Before call to gets

Stack Frame for main

Return Address

Saved %ebp

[3][2][1][0] buf

Stack Frame for echo

```

/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}

```

```

echo:
    pushl %ebp          # Save %ebp on stack
    movl %esp, %ebp
    pushl %ebx          # Save %ebx
    leal -8(%ebp),%ebx  # Compute buf as %ebp-8
    subl $20, %esp      # Allocate stack space
    movl %ebx, (%esp)   # Push buf on stack
    call gets           # Call gets
    . . .

```

System 1

Buffer Overflow Stack Example

```

unix> gdb bufdemo
(gdb) break echo
Breakpoint 1 at 0x8048583
(gdb) run
Breakpoint 1, 0x8048583 in echo ()
(gdb) print /x %ebp
$1 = 0xffffc638
(gdb) print /x *(unsigned *)$ebp
$2 = 0xffffc658
(gdb) print /x *((unsigned *)$ebp + 1)
$3 = 0x80485f7

```

Before call to gets

Stack Frame for main

Return Address

Saved %ebp

[3][2][1][0] buf

Stack Frame for echo

Stack Frame for main

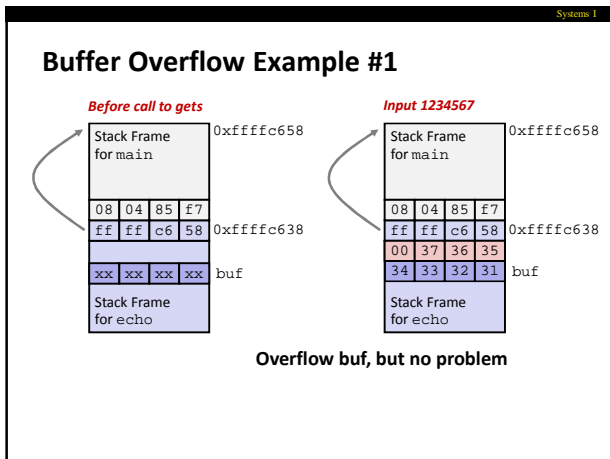
08 04 85 f7

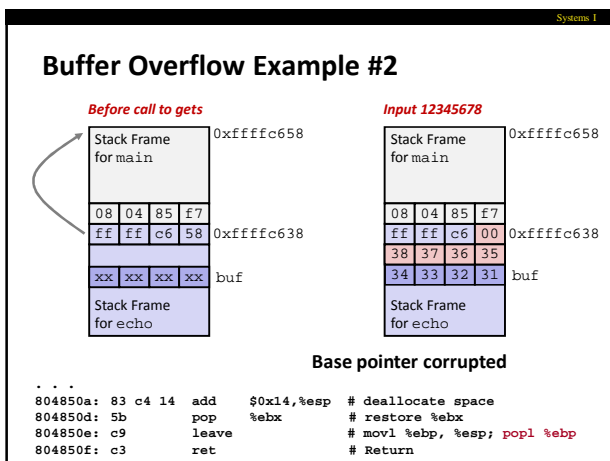
ff ff c6 58

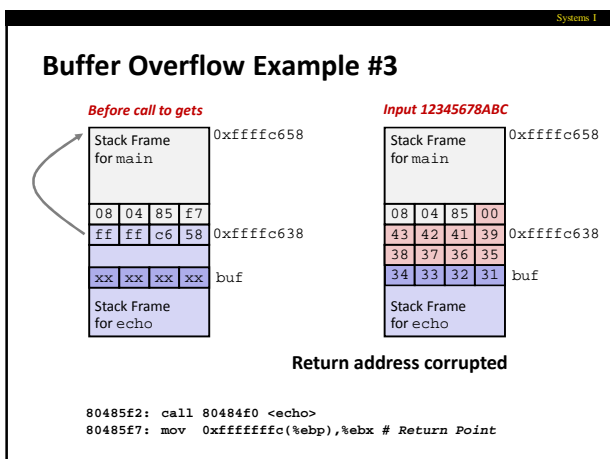
xx xx xx xx

Stack Frame for echo

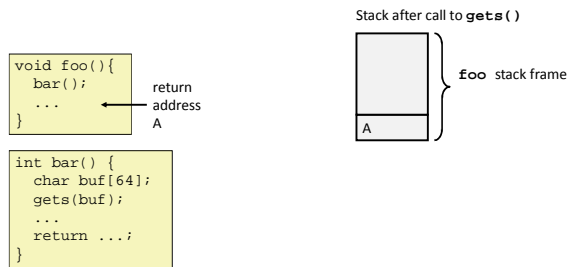
80485f2: call 80484f0 <echo>
80485f7: mov 0xffffffff(%ebp),%ebx # Return Point



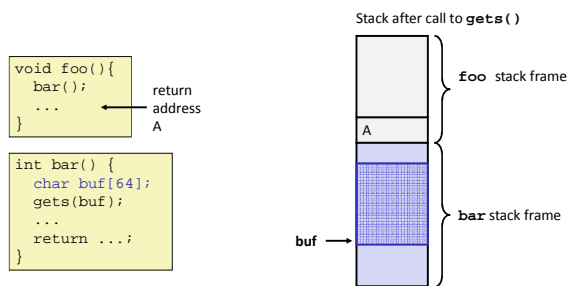




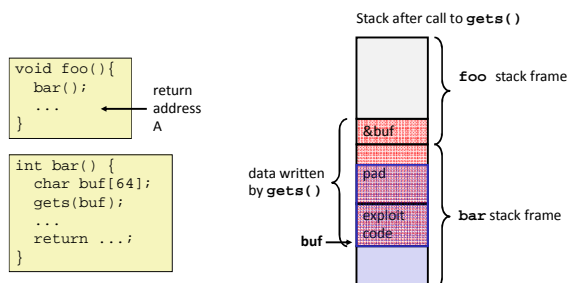
Malicious Use of Buffer Overflow



Malicious Use of Buffer Overflow



Malicious Use of Buffer Overflow



- Input string contains byte representation of executable code
- Overwrite return address with address of buffer
- When `bar()` executes `ret`, will jump to exploit code

Exploits Based on Buffer Overflows

- *Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines*
- **Internet worm**
 - Early versions of the finger server (fingerd) used `gets()` to read the argument sent by the client:
 - `finger droh@cs.cmu.edu`
 - Worm attacked fingerd server by sending phony argument:
 - `finger "exploit-code padding new-return-address"`
 - exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.

Exploits Based on Buffer Overflows

- *Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines*
- **IM War**
 - AOL exploited existing buffer overflow bug in AIM clients
 - exploit code: returned 4-byte signature (the bytes at some location in the AIM client) to server.
 - When Microsoft changed code to match signature, AOL changed signature location.

Date: Wed, 11 Aug 1999 11:30:57 -0700 (PDT)
 From: Phil Bucking <philbucking@yahoo.com>
 Subject: AOL exploiting buffer overrun bug in their own software!
 To: rms@pharlap.com

Mr. Smith,

I am writing you because I have discovered something that I think you might find interesting because you are an Internet security expert with experience in this area. I have also tried to contact AOL but received no response.

I am a developer who has been working on a revolutionary new instant messaging client that should be released later this year.

...
 It appears that the AIM client has a buffer overrun bug. By itself this might not be the end of the world, as MS surely has had its share. But AOL is now *exploiting their own buffer overrun bug* to help in its efforts to block MS Instant Messenger.

....
 Since you have significant credibility with the press I hope that you can use this information to help inform people that behind AOL's friendly exterior they are nefariously compromising peoples' security.

Sincerely,
 Phil Bucking
 Founder, Bucking Consulting
 philbucking@yahoo.com


It was later determined that this email originated from within Microsoft!

Code Red Worm

- [illegible]

Code Red Exploit Code

- Starts 100 threads running
- Spread self
 - Generate random IP addresses & send attack string
 - Between 1st & 19th of month
- Attack www.whitehouse.gov
 - Send 98,304 packets; sleep for 4-1/2 hours; repeat
 - Denial of service attack
 - Between 21st & 27th of month
- Deface server's home page
 - After waiting 2 hours



Code Red Effects

- **Later Version Even More Malicious**
 - Code Red II
 - As of April, 2002, over 18,000 machines were infected
 - Was still spreading
- **Paved Way for NIMDA**
 - Variety of propagation methods
 - One was to exploit vulnerabilities left behind by Code Red II
- **ASIDE (security flaws start at home)**
 - .rhosts used by Internet Worm
 - Attachments used by MyDoom

Avoiding Overflow Vulnerability

```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    fgets(buf, 4, stdin);
    puts(buf);
}
```

- Use library routines that limit string lengths
 - `fgets` instead of `gets`
 - `strncpy` instead of `strcpy`
 - Don't use `scanf` with `%s` conversion specification
 - Use `fgets` to read the string
 - Or use `%ns` where `n` is a suitable integer

System-Level Protections

- Randomized stack offsets
 - At start of program, allocate random amount of space on stack
 - Makes it difficult for hacker to predict beginning of inserted code
- Nonexecutable code segments
 - In traditional x86, can mark region of memory as either "read-only" or "writeable"
 - Can execute anything readable
 - Add explicit "execute" permission

```
unix> gdb bufdemo
(gdb) break echo

(gdb) run
(gdb) print /x $ebp
$1 = 0xffffc638

(gdb) run
(gdb) print /x $ebp
$2 = 0xffffbb08

(gdb) run
(gdb) print /x $ebp
$3 = 0xffffc6a8
```

Worms and Viruses

- Worm: A program that
 - Can run by itself
 - Can propagate a fully working version of itself to other computers
- Virus: Code that
 - Add itself to other programs
 - Cannot run independently
- Both are (usually) designed to spread among computers and to wreak havoc
