

Computer Systems I

Class 9

Alignment

- **Aligned Data**
 - Primitive data type requires K bytes
 - Address must be multiple of K
 - Required on some machines; advised on IA32
 - treated differently by IA32 Linux, x86-64 Linux, and Windows!
- **Motivation for Aligning Data**
 - Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - Inefficient to load or store datum that spans quad word boundaries
 - Virtual memory very tricky when datum spans 2 pages
- **Compiler**
 - Inserts gaps in structure to ensure correct alignment of fields

Specific Cases of Alignment (IA32)

- **1 byte: `char`, ...**
 - no restrictions on address
- **2 bytes: `short`, ...**
 - lowest 1 bit of address must be 0₂
- **4 bytes: `int`, `float`, `char *`, ...**
 - lowest 2 bits of address must be 00₂
- **8 bytes: `double`, ...**
 - Windows (and most other OS's & instruction sets):
 - lowest 3 bits of address must be 000₂
 - Linux:
 - lowest 2 bits of address must be 00₂
 - i.e., treated the same as a 4-byte primitive data type


Specific Cases of Alignment (x86-64)

- **1 byte: char, ...**
 - no restrictions on address
- **2 bytes: short, ...**
 - lowest 1 bit of address must be 0_2
- **4 bytes: int, float, ...**
 - lowest 2 bits of address must be 00_2
- **8 bytes: double, char *, ...**
 - Windows & Linux:
 - lowest 3 bits of address must be 000_2

Satisfying Alignment with Structures

- **Within structure:**
 - Must satisfy element's alignment requirement
- **Overall structure placement**
 - Each structure has alignment requirement K
 - K = Largest alignment of any element
 - Initial address & structure length must be multiples of K
- **Example (under Windows or x86-64):**
 - K = 8, due to **double** element


```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```




Different Alignment Conventions

- **x86-64 or IA32 Windows:**
 - K = 8, due to **double** element

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```



- **IA32 Linux**
 - K = 4; **double** treated like a 4-byte data type



Saving Space


- Put large data types first

```
struct S1 {
  char c;
  int i[2];
  double v;
} *p;
```



```
struct S2 {
  double v;
  int i[2];
  char c;
} *p;
```

- Effect (example x86-64, both have $K=8$)




Arrays of Structures

- Padding has been added inside a struct so that fields are aligned.
- This only works if the entire struct is aligned (starts at a correct address).
- What effect does this have on arrays of structs?

Arrays of Structures

- Satisfy alignment requirement for every element

```
struct S2 {
  double v;
  int i[2];
  char c;
} a[10];
```



Accessing Array Elements

- Compute array offset 12i
- Compute offset 8 with structure
- Assembler gives offset a+8
 - Resolved during linking

```

struct S3 {
    short s1;
    float v;
    short s2;
} a[10];

```

```

short get_s2(int idx)
{
    return a[idx].s2;
}

```

```

# %eax = idx
leal (%eax,%eax,2),%eax # 3*idx
movswl a+8(,%eax,4),%eax

```

Unions

- What is a union?
- How does it differ from a struct?

Union Allocation

- Allocate according to largest element
- Can only use ones field at a time

```

union U1 {
    char c;
    int i[2];
    double v;
} *up;

```

```

struct S1 {
    char c;
    int i[2];
    double v;
} *sp;

```

Using Union to Access Bit Patterns

```
typedef union {
    float f;
    unsigned u;
} bit_float_t;
```



Don't forget byte ordering!

```
float bit2float(unsigned u)
{
    bit_float_t arg;
    arg.u = u;
    return arg.f;
}
```

Same as (float) u ?

```
unsigned float2bit(float f)
{
    bit_float_t arg;
    arg.f = f;
    return arg.u;
}
```

Same as (unsigned) f ?

Byte Ordering Revisited

- **Idea**
 - Short/long/quad words stored in memory as 2/4/8 consecutive bytes
 - Which is most (least) significant?
 - Can cause problems when exchanging binary data between machines
- **Big Endian**
 - Most significant byte has lowest address
 - PowerPC, Sparc
- **Little Endian**
 - Least significant byte has lowest address
 - Intel x86

Byte Ordering Example

```
union {
    unsigned char c[8];
    unsigned short s[4];
    unsigned int i[2];
    unsigned long l[1];
} dw;
```

c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
s[0]	s[1]	s[2]	s[3]				
i[0]		i[1]					
l[0]							

Byte Ordering Example (Cont).

```
int j;
for (j = 0; j < 8; j++)
    dw.c[j] = 0xf0 + j;

printf("Characters 0-7 ==
[0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x]\n",
    dw.c[0], dw.c[1], dw.c[2], dw.c[3],
    dw.c[4], dw.c[5], dw.c[6], dw.c[7]);


printf("Shorts 0-3 ==
[0x%x,0x%x,0x%x,0x%x]\n",
    dw.s[0], dw.s[1], dw.s[2], dw.s[3]);

printf("Ints 0-1 == [0x%x,0x%x]\n",
    dw.i[0], dw.i[1]);

printf("Long 0 == [0x%lx]\n",
    dw.l[0]);
```

Byte Ordering on IA32

Little Endian




Output on IA32:

```
Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]
Shorts     0-3 == [0xf1f0,0xf3f2,0xf5f4,0xf7f6]
Ints       0-1 == [0xf3f2f1f0,0xf7f6f5f4]
Long       0 == [0xf3f2f1f0]
```

Byte Ordering on Sun

Big Endian



Output on Sun:

```
Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]
Shorts     0-3 == [0xf0f1,0xf2f3,0xf4f5,0xf6f7]
Ints       0-1 == [0xf0f1f2f3,0xf4f5f6f7]
Long       0 == [0xf0f1f2f3]
```

Systeme I

Byte Ordering on x86-64

Little Endian

Output on x86-64:

Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]
Shorts 0-3 == [0xf1f0,0xf3f2,0xf5f4,0xf7f6]
Ints 0-1 == [0xf3f2f1f0,0xf7f6f5f4]
Long 0 == [0xf7f6f5f4f3f2f1f0]

Systeme I

Summary

- **Arrays in C**
 - Contiguous allocation of memory
 - Aligned to satisfy every element's alignment requirement
 - Pointer to first element
 - No bounds checking
- **Structures**
 - Allocate bytes in order declared
 - Pad in middle and at end to satisfy alignment
- **Unions**
 - Overlay declarations
 - Way to circumvent type system

Systeme I

Struct review

```
typedef struct {
    double one;
    char two[5];
    int three;
    int *four;
    char five;
} struct_one;

struct_one array[10];

void main() {
    printf("address of array = %d\n", array);
    printf("address of array[3] = %d\n", &array[3]);
    printf("address of array[5].two[2] = %d\n", &array[5].two[2]);
}
```

address of array = 10000

Union review

```
void main() {
    union_one U;
    U.number[0] = 0x01234567;
    U.number[1] = 0x89abcdef;
    printf("number[0] = %08x\n", U.number[0]);
    printf("number[1] = %08x\n", U.number[1]);
    printf("character[0-3] = %02x%02x%02x%02x\n",
        U.character[0], U.character[1], U.character[2], U.character[3]);
    printf("character[4-7] = %02x%02x%02x%02x\n",
        U.character[4], U.character[5], U.character[6], U.character[7]);
}
```

number[0] = ???
 number[1] = ???
 character[0-3] = ???
 character[4-7] = ???

Summary for last week

- **Arrays in C**
 - Contiguous allocation of memory
 - Aligned to satisfy every element's alignment requirement
 - Pointer to first element
 - No bounds checking
- **Structures**
 - Allocate bytes in order declared
 - Pad in middle and at end to satisfy alignment
- **Unions**
 - Overlay declarations
 - Way to circumvent type system
